

Shiv Nadar University Chennai

CS3807 – Deep Learning Laboratory

Experiment 5

Comprehensive Study of CNN Training, Regularization, Optimization, Hyperparameter Tuning, Transfer Learning and Cross-Validation

Degree & Branch	B.Tech Artificial Intelligence & Data Science	Semester V
Subject Code & Name	CS3807 – Deep Learning Laboratory	AY: 2026–27

Name : Ashwin S
Reg no : 24011101007
Class : AI DS - A (batch 1)

1. Objective

The objective of this experiment is to systematically study the effect of weight initialization, regularization, optimization algorithms, CNN hyperparameters, transfer learning, fine-tuning and cross-validation on image classification performance.

The experiment uses a single CNN architecture, MobileNetV2, and the Oxford-IIIT Pet dataset. Students will analyze the effect of individual design choices using training/validation curves and finally select a suitable configuration using 5-fold cross-validation.

2. Learning Outcomes

After completing this experiment, students will be able to:

- explain different weight initialization strategies;
- identify and analyze overfitting using training and validation curves;
- compare SGD, Momentum, RMSProp and Adam;
- explain CNN hyperparameters such as kernel size, stride, padding and pooling;
- understand Batch Normalization and Dropout;
- perform transfer learning and fine-tuning using MobileNetV2;
- tune important training hyperparameters;
- apply 5-fold cross-validation for model selection; and
- justify the final model using accuracy, variability and computational cost.

3. Dataset and Experimental Setup

Dataset: Oxford-IIIT Pet Dataset

The Oxford-IIIT Pet Dataset contains images of cats and dogs belonging to 37 breeds. The images are RGB and have different spatial dimensions. For this experiment, all images are resized to

$$224 \times 224 \times 3.$$

The dataset is suitable for demonstrating transfer learning using ImageNet-pretrained CNNs.

Experimental considerations:

- Use RGB images resized to 224×224 .
- Normalize the input images according to the pretrained model requirements.
- Maintain separate training, validation and test data.
- The test set must remain untouched during hyperparameter selection.
- Use MobileNetV2 because it is computationally lighter than many large CNN architectures and is more suitable for CPU-based laboratory work.


4. MobileNetV2 Architecture

MobileNetV2 is a lightweight CNN architecture designed to reduce computational complexity while maintaining good image representation.

Its important components include:

- depthwise convolution;
- pointwise 1×1 convolution;
- inverted residual blocks;
- linear bottlenecks;
- batch normalization; and
- ReLU6 activation.

A simplified representation is:



Note: The above diagram is a simplified conceptual representation rather than the complete MobileNetV2 architecture.

5. Weight Initialization

Weight initialization determines the initial values of trainable weights before training begins. A suitable initialization can improve convergence and reduce problems such as vanishing or exploding activations.

Four schemes were investigated on the newly added classifier head, with the pretrained MobileNetV2 base kept frozen: Zero initialization, Random (normal) initialization, Xavier/Glorot initialization, and He initialization.




Figure 1: Training Loss vs. Epoch for different weight initialization schemes.

Zero initialization stays flat at a loss of 3.611 across all eight epochs, since every neuron gets the same gradient update and none of them can break symmetry. With same starting weights, every neuron produces the same output and gradient forever. Random, Xavier and He initialization all fall from around 1.2–1.3 down to 0.17–0.20, and have very similar results to each other.




Figure 2: Validation Accuracy vs. Epoch for different weight initialization schemes.

Zero initialization stays near random-guess accuracy throughout, while Random, Xavier and He all climb quickly and plateau at a high accuracy within the first couple of epochs. This happens because the frozen MobileNetV2 base already supplies strong features, so any non-zero initialization lets the head learn almost immediately.

Xavier/Glorot and He gave the fastest, most stable convergence among the non-zero schemes, while Zero initialization never trained at all.

6. Regularization and Overfitting

Overfitting occurs when a model performs very well on the training data but generalizes poorly to unseen data. Four configurations were compared: No regularization, L2 regularization, Dropout, and Batch Normalization.




Figure 3: Training and Validation Accuracy vs. Epoch for each regularization setting.

No Regularization and Batch Normalization both push training accuracy close to 100% while validation accuracy plateaus much lower, leaving a clear gap. Dropout keeps the two curves closest together, since it stops the head from memorizing individual training images the way the other two settings allow.




Figure 4: Training and Validation Loss vs. Epoch for each regularization setting.

For No Regularization and Batch Normalization, training loss keeps falling while validation loss flattens out and drifts upward again after a few epochs, a classic overfitting sign. Dropout is the only setting where validation loss keeps falling alongside training loss for almost the whole run.

7. Batch Normalization

Batch Normalization normalizes activations within a mini-batch during training.

For a mini-batch

$$x_1, x_2, \dots, x_m,$$

the batch mean is

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$$

and the batch variance is

$$\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2.$$

The normalized activation is

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}.$$

The final output is

$$y_i = \gamma \hat{x}_i + \beta,$$

where γ and β are learnable parameters.

Numerical Example

Consider

$$x = [2, 4, 6, 8].$$

The mean is

$$\mu_B = \frac{2 + 4 + 6 + 8}{4} = 5.$$

The variance is

$$\sigma_B^2 = \frac{(2-5)^2 + (4-5)^2 + (6-5)^2 + (8-5)^2}{4} = 5.$$

Therefore,

$$\sqrt{5} \approx 2.236.$$

Ignoring the very small ϵ ,

$$\hat{x}_1 = \frac{2-5}{2.236} \approx -1.342,$$

$$\hat{x}_2 = \frac{4-5}{2.236} \approx -0.447,$$

$$\hat{x}_3 = \frac{6-5}{2.236} \approx 0.447,$$

$$\hat{x}_4 = \frac{8-5}{2.236} \approx 1.342.$$

Hence,

$$\hat{x} \approx [-1.342, -0.447, 0.447, 1.342]$$

If $\gamma = 1$ and $\beta = 0$, these are the final outputs.

Inference: Batch Normalization produces approximately zero-centred and unit-scaled activations for this mini-batch. The learnable γ and β allow the network to learn an appropriate scale and shift.




Figure 5: Validation Accuracy vs. Epoch, With BN vs. Without BN.

The two curves stay close together for most of training and reach similar peak accuracy. Batch Normalization stabilizes the activation distribution feeding the head, which can smooth out fluctuations, but the effect is modest here since the base network itself is frozen.

8. Optimization Algorithms

Four optimization methods were compared: SGD, Momentum, RMSProp and Adam.




Figure 6: Training Loss vs. Epoch for different optimizers.

SGD barely moves off its starting loss, while Momentum, RMSProp and Adam all drop sharply within the first few epochs. Adam and RMSProp reach the lowest final loss because they adapt their step size per parameter, letting them take large useful steps early on.




Figure 7: Validation Accuracy vs. Epoch for different optimizers.

Adam converges fastest, RMSProp is close behind, Momentum trails both, and SGD stays far below the other three throughout. Plain SGD has no adaptive step size, so it needs many more epochs than were budgeted here to catch up.

Optimizer	Final Loss	Best Val. Accuracy	Epoch to Converge	Time
SGD	01.280	79.96%	10	134.200 s
Momentum	00.355	89.17%	9	125.100 s
RMSProp	00.155	89.71%	10	125.000 s
Adam	00.141	89.98%	9	125.400 s

9. CNN Hyperparameter Tuning

The following hyperparameters were investigated, changing one at a time while keeping the remaining settings fixed at their default values (Learning Rate = 0.001, Batch Size = 32, Dropout Rate = 0.25, Optimizer = Adam):

Hyperparameter	Suggested Values
Learning Rate	0.001, 0.0001
Batch Size	16, 32, 64
Dropout Rate	0, 0.25, 0.5
Optimizer	SGD, Adam
Fine-Tuning Learning Rate	10^{-4} , 10^{-5}
Frozen Layers	Frozen base / Partial unfreezing

For CNN concepts, students should also understand:

- kernel/filter;
- stride;
- padding;
- pooling;
- number of channels; and
- depthwise separable convolution.

For a convolution operation, the output dimension can be calculated as

$$O = \left\lfloor \frac{N + 2P - K}{S} \right\rfloor + 1,$$

where N is input size, K is kernel size, P is padding and S is stride. For $N = 224$, $K = 3$, $P = 1$, $S = 1$, this gives $O = 224.000$, and for $N = 224$, $K = 3$, $P = 0$, $S = 2$, this gives $O = 111.000$.




Figure 8: Learning Rate vs. Validation Accuracy.

The larger learning rate reaches higher validation accuracy than the smaller one within this epoch budget. A smaller learning rate moves the weights a shorter distance each step, so it simply has not converged yet by the time training stops.




Figure 9: Batch Size vs. Validation Accuracy.

The middle batch size gives the best result, with the two extremes close behind and no strictly monotonic trend. This is expected since batch size trades off gradient noise against update frequency in both directions.




Figure 10: Dropout Rate vs. Validation Accuracy.

A moderate dropout rate performs best, while both very low and very high dropout perform worse, forming a U-shape. Too little dropout lets the head overfit, while too much removes so much information that it struggles to learn at all.

Based on these three sweeps, the best hyperparameters found were a learning rate of 0.001, a batch size of 32, and a dropout rate of 0.25.

10. Transfer Learning and Fine-Tuning

MobileNetV2 pretrained on ImageNet was used.

Case A: Feature Extraction

Pretrained MobileNetV2 → Freeze Base → New Classifier

The pretrained convolutional layers are frozen and only the newly added classification layers are trained.

Case B: Fine-Tuning

Pretrained MobileNetV2 → Unfreeze Selected Layers → Small Learning Rate

The upper layers of the pretrained network are unfrozen and trained with a smaller learning rate.




Figure 11: Feature Extraction vs. Fine-Tuning: Validation Accuracy vs. Epoch.

Feature extraction climbs steadily on its own, and once fine-tuning begins, validation accuracy pushes a little higher still. Unfreezing the top layers lets the network adapt its higher-level filters to pet breeds instead of relying only on generic ImageNet features.




Figure 12: Training and Validation Loss across Feature Extraction and Fine-Tuning.

Training loss keeps falling smoothly across both phases, while validation loss stays roughly flat rather than rising sharply. This suggests the small fine-tuning learning rate was gentle enough to keep improving the model without erasing the pretrained features.

Fine-tuning normally uses a smaller learning rate than a freshly initialized classifier because the pretrained weights already hold useful features learned from a much larger dataset. A large learning rate would overwrite this knowledge almost immediately, so a small one adapts it gently instead.

11. K-Fold Cross-Validation

After the individual studies above, four promising configurations were selected for 5-fold cross-validation on the training data: C1 (Adam, dropout 0.5, Batch Normalization), C2 (Adam, low learning rate 10^{-4} , dropout 0.25), C3 (RMSProp, dropout 0.5, Batch Normalization), and C4 (Adam with L2 regularization).

For $K = 5$,

$$\bar{A} = \frac{1}{5} \sum_{i=1}^5 A_i$$

and

$$SD = \sqrt{\frac{1}{5} \sum_{i=1}^5 (A_i - \bar{A})^2}.$$

Report the result as

$$\boxed{\bar{A} \pm SD}.$$

The independent test set was not used during this hyperparameter selection stage.

Configuration	F1	F2	F3	F4	F5	Mean $\pm$ SD
C1: Adam, Dropout 0.5, BN	91.69%	91.50%	90.05%	91.68%	92.65%	91.51% $\pm$ 00.840
C2: Adam, Low LR, Dropout 0.25	89.76%	87.25%	87.05%	89.46%	88.88%	88.48% $\pm$ 01.120
C3: RMSProp, Dropout 0.5, BN	91.79%	91.11%	91.01%	91.68%	92.55%	91.63% $\pm$ 00.550
C4: Adam, L2 Regularization	91.59%	90.72%	89.18%	90.72%	93.13%	91.07% $\pm$ 01.290




Figure 13: 5-Fold Cross-Validation Accuracy with standard-deviation error bars.

C3 has both the highest mean accuracy and the smallest error bar, making it the most reliable choice. C2 has the lowest mean and the widest spread. Averaging over several folds shows up configurations that are consistent, not just ones that got lucky on a single split.

Based on both mean accuracy and standard deviation across folds, C3 was selected as the final configuration for retraining and test-set evaluation.

12. Final Model Evaluation

The selected configuration, C3 (RMSProp, dropout 0.5, Batch Normalization), was retrained using the complete training data for 12 epochs, with the independent test set held out until this stage.

Metric	Value
Mean CV Accuracy	91.630%
CV Standard Deviation	00.550
Test Accuracy	90.800%
Precision	00.912
Recall	00.908
F1-score	00.907
Training Time	143.300 s
Number of Parameters	2431845




Figure 14: Confusion Matrix of the final model on the test set.

The model reaches 100% accuracy on several visually distinctive breeds such as Siamese and Newfoundland. The most confused pairs are breeds that look alike up close, such as American Pit Bull Terrier with Staffordshire Bull Terrier, and Basset Hound with Beagle. Almost all confusion happens within species and between breeds sharing coat length, colouring or body shape, details that are easy to lose after downsampling and pooling.

Plot 15 (Optional): Representative Misclassified Images




Figure 15: Representative misclassified images from the test set.

13. Overall Results

Configuration	Test/Validation Accuracy
Baseline (No Regularization)	90.430% (val)
Best Initialization (Random)	90.250% (val)
Best Regularization (No Regularization)	90.430% (val)
Best Optimizer (Adam)	89.980% (val)
Best Hyperparameters (from sweeps)	90.700% (val)
Fine-Tuned Model	90.970% (val)
Final Selected Model (C3: RMSProp, Dropout 0.5, BN)	90.800% (test)

For the final selected model specifically: Mean CV Accuracy = 91.630%, CV Standard Deviation = 00.550, Training Time = 143.300 s.

The final selected model was chosen even though the fine-tuned model reaches a slightly higher single-run validation accuracy, because C3 was the only configuration evaluated with 5-fold cross-validation. Its low standard deviation across folds shows the result is unlikely to swing much on a different split, and its test accuracy landed close to its CV mean, which together make it the more trustworthy pick.

14. Source Code

The complete implementation, including data preprocessing, weight initialization, regularization, batch normalization, optimization algorithms, hyperparameter tuning, transfer learning, fine tuning, evaluation, and the additional exercises, is available in the following repository :

<https://github.com/Ashwin271206/dl-lab/tree/main/Lab%205>

The repository contains a single Kaggle notebook covering all tasks experimented in this report, along with the generated figures and outputs.

15. Discussion

1. What is the difference between model parameters and hyperparameters?

Model parameters (weights and biases) are values learned automatically during training through gradient descent. Hyperparameters (learning rate, batch size, dropout rate, number of epochs, etc) are set before training begins and control how the learning process itself behaves. Parameters are learned from data and hyperparameters are given by the user.

2. Why is weight initialization important?

Weight initialization determines how fast or how reliably the model converges. Poor initialization can cause vanishing or exploding gradients in deep networks, slow convergence, or, in the worst case, prevent the model from learning at all.

3. Why can zero initialization be problematic for neural networks?

If every weight in a layer starts at zero, every neuron in that layer computes the exact same output and receives the exact same gradient during backpropagation. The neurons never differentiate from one another, so the layer behaves as if it had only a single neuron regardless of its actual width.

4. Compare Xavier and He initialization.

Both scale the initial weight variance based on the number of input units to a layer, in order to keep the variance of activations roughly stable across layers. Xavier/Glorot initialization is derived assuming a linear or symmetric activation function such as tanh or sigmoid, while He initialization uses a larger variance derived for ReLU-family activations, which zeroes out roughly half their inputs and would otherwise shrink activation variance layer by layer if Xavier's scaling were used instead.

5. How can training and validation curves be used to identify overfitting?

Overfitting can be noticed through a growing gap between training and validation performance, when training accuracy keeps rising (or training loss keeps falling) while validation accuracy plateaus or falls (or validation loss stops falling and starts rising).

6. How does Dropout reduce overfitting?

Dropout randomly deactivates a fraction of neurons on each training step, which prevents the network from relying too heavily on any single unit or a small group of units. This forces the remaining active units to learn more redundant, generally useful features.

7. What is the purpose of Batch Normalization?

Batch Normalization normalizes the activations passing through a layer to have approximately zero mean and unit variance within each mini-batch, then rescales and shifts them using learnable parameters. This stabilizes the distribution of inputs each layer sees during training, which can speed up convergence and reduce sensitivity to the initial weight scale.

8. Explain the numerical Batch Normalization example.

For the mini-batch $x = [2, 4, 6, 8]$, the mean is $\mu_B = 5$ and the variance is $\sigma_B^2 = 5$. Subtracting the mean and dividing by the standard deviation ($\sqrt{5} \approx 2.236$) gives the normalized values $\hat{x} \approx [-1.342, -0.447, 0.447, 1.342]$, which have approximately zero mean and unit variance. With $\gamma = 1$ and $\beta = 0$, these normalized values pass through unchanged as the final output of the batch normalization layer.

9. What are the roles of γ and β ?

After normalization, γ and β are learnable parameters that let the model scale and shift the normalized activations respectively. This matters because forcing every activation to have exactly zero mean and unit variance can

be an unnecessarily strict constraint; γ and β let the network recover a different scale or shift if that turns out to work better for a particular layer.

10. Compare SGD, Momentum, RMSProp and Adam.

SGD updates weights using only the current gradient scaled by a fixed learning rate, which made it converge far slower than the other three optimizers. Momentum accumulates a moving average of past gradients to smooth out updates and accelerate progress along consistent directions. RMSProp divides each parameter's update by a running average of recent squared gradients, giving each parameter its own adaptive step size. Adam combines both ideas, tracking both a momentum term and a per-parameter adaptive scale, which is why it reached the best validation accuracy and lowest final loss among the four optimizers tested here.

11. What happens when the learning rate is too large?

Updates overshoot the minimum of the loss surface, which can cause training loss to oscillate, diverge, or fail to decrease consistently even though the optimizer is, in principle, moving toward a good solution.

12. What happens when the learning rate is too small?

Training becomes very slow, since each update only moves the weights a tiny distance. Given a fixed number of epochs, as used throughout this experiment, a learning rate that is too small may fail to converge at all within the available budget, which matches the weaker result seen for the 0.0001 learning rate in Section 9's sweep compared with 0.001.

13. What is the effect of increasing batch size?

Larger batches give a smoother, lower-variance estimate of the gradient at each step, but produce fewer weight updates per epoch. Smaller batches update more frequently and introduce more noise into training, which can help generalization but also makes each individual gradient step less reliable.

14. Explain stride and padding.

Stride is how much the filter moves across the input image after each computation. A stride of 1 slides the filter one pixel at a time, while a stride of 2 skips every alternate pixel and so on. Higher the stride, lower the output image's size. Padding adds extra pixels, generally 0s, around the border of the input before convolution process, which maintains the size of the output same as input size.

15. Why is MobileNetV2 computationally efficient?

MobileNetV2 replaces standard convolutions with depthwise separable convolutions, which factor a normal convolution into a per-channel spatial convolution followed by a 1×1 convolution that mixes channels. This factorization needs far fewer multiply-accumulate operations and parameters than a standard convolution of the same input/output size, which is also why the final model in this experiment has only about 2.432 million parameters despite handling 37-class classification on 224×224 images.

16. What is depthwise separable convolution?

It is a two-step convolution: first, a depthwise convolution applies a single filter independently to each input channel, capturing spatial patterns per channel without mixing channels together; second, a pointwise (1×1) convolution combines the outputs of all channels to produce the final feature maps. Splitting the operation this way achieves a similar receptive field and expressive power to a standard convolution at a fraction of the computational cost.

17. What is transfer learning?

Transfer learning reuses a model (or part of one) that has already been trained on a large dataset, such as ImageNet, as a starting point for a new, often smaller, dataset and task. Rather than learning visual features such as edges, textures and shapes from scratch, the new model starts with these already learned and only needs to adapt them to the new task, which is exactly the approach used throughout this experiment with pretrained MobileNetV2.

18. Differentiate feature extraction and fine-tuning.

In feature extraction, the pretrained convolutional base is entirely frozen and only a newly added classifier head is trained, treating the base purely as a fixed feature generator. In fine-tuning, some of the upper layers of the pretrained base are unfrozen and trained alongside the classifier head, usually at a smaller learning rate, letting the network adapt its higher-level features to the new dataset.

19. Why is a smaller learning rate generally used during fine-tuning?

The unfrozen pretrained layers already have well-tuned weights learned from a much larger dataset. A large learning rate would produce large updates that could overwrite this pretrained knowledge within just a few steps,

and cause catastrophic forgetting. A small learning rate makes small changes that adapt the pretrained features to the new task without destroying what they already encode.

20. Why is K-Fold Cross-Validation useful for hyperparameter selection?

A single train/validation split can give a misleading estimate of how well a configuration generalizes. K-fold cross-validation trains and evaluates the same configuration across K different splits of the data, producing both a mean accuracy and a standard deviation that together describe generalization.

21. Why must the test set remain untouched during tuning?

If the test set influences any part during training or hyperparameter selection, even indirectly, the final test accuracy would no longer show how the model would perform on genuinely unseen data. Keeping the test set completely separate until evaluation ensures that the results are based on totally unseen data.

22. Why should mean and standard deviation both be reported?

The mean alone describes typical performance but says nothing about consistency. Two configurations can have nearly identical mean accuracy while one is far more stable across different data splits than the other, as seen when comparing C3's standard deviation of 00.550 against C4's 01.290 in Section 11, even though their means differed by less than one percentage. Hence, standard deviation is also important.

23. Is the highest validation accuracy always sufficient to select a model?

No. A single validation accuracy number can be optimistic by chance, especially with a small validation set, and says nothing about how consistent that performance would be across different data splits. In this experiment, the fine-tuned model reached a higher single-run validation accuracy (90.97%) than the cross-validated C3 configuration's mean (91.63% CV mean, 90.80% test), yet C3 was preferred for the final model because its performance had been verified across multiple folds with low variability rather than a single good split.

16. Additional Exercise

Two new combinations of learning rate, dropout, batch size and fine-tuning strategy were evaluated using 5-fold cross-validation and compared against the previously selected configuration, C3:

- **N1:** Adam, learning rate 10^{-4} , dropout 0.5, batch size 16, with fine-tuning of the last 20 base layers at a reduced learning rate.
- **N2:** Adam, learning rate 10^{-3} , dropout 0.1, batch size 64, feature extraction only (no fine-tuning).

Configuration	F1	F2	F3	F4	F5	Mean $\pm$ SD
N1: Low LR, High Dropout, Small Batch, FT	91.40%	90.14%	89.66%	90.43%	91.59%	90.64% $\pm$ 00.740
N2: High LR, Low Dropout, Large Batch, FE	90.72%	91.21%	90.14%	90.91%	92.55%	91.11% $\pm$ 00.800
C3: RMSProp, Dropout 0.5, BN (selected)	91.79%	91.11%	91.01%	91.68%	92.55%	91.63% $\pm$ 00.550




Figure 16: New configurations (N1, N2) compared against the previously selected best configuration (C3) via 5-Fold CV.

Neither new configuration beats C3. N2 comes closer but still falls short on both mean accuracy and standard deviation, while N1 trails both. C3 stays the better pick on every count asked for: higher mean accuracy, lower spread across folds, and a computational cost already shown to be reasonable in Section 12.

17. Conclusion

This experiment performed the full deep learning workflow for image classification, using MobileNetV2 on the Oxford-IIIT Pet dataset’s 37 cat and dog breeds. It covered weight initialization, regularization, Batch Normalization, optimization algorithms, hyperparameter tuning, transfer learning, fine-tuning, and cross-validation.

A pretrained MobileNetV2 was used with the base frozen at first and a new classifier head trained on top. Zero, Random, Xavier and He initialization were compared, along with L2 regularization, Dropout and Batch Normalization. SGD, Momentum, RMSProp and Adam were tested as optimizers, followed by sweeps over learning rate, batch size and dropout rate. Feature extraction was then compared against fine-tuning the last layers of MobileNetV2, and several configurations were run through 5-fold cross-validation to pick a final model based on both accuracy and consistency, before a last check on the untouched test set.

18. References

1. Ian Goodfellow, Yoshua Bengio and Aaron Courville, *Deep Learning*, MIT Press, 2016.
2. Sergey Ioffe and Christian Szegedy, “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift,” ICML, 2015.
3. Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Zhmoginov and Liang-Chieh Chen, “MobileNetV2: Inverted Residuals and Linear Bottlenecks,” CVPR, 2018.
4. Omkar M. Parkhi, Andrea Vedaldi, Andrew Zisserman and C. V. Jawahar, “Cats and Dogs,” IEEE Conference on Computer Vision and Pattern Recognition, 2012.
5. TensorFlow Documentation: <https://www.tensorflow.org>
6. Keras Documentation: <https://keras.io>